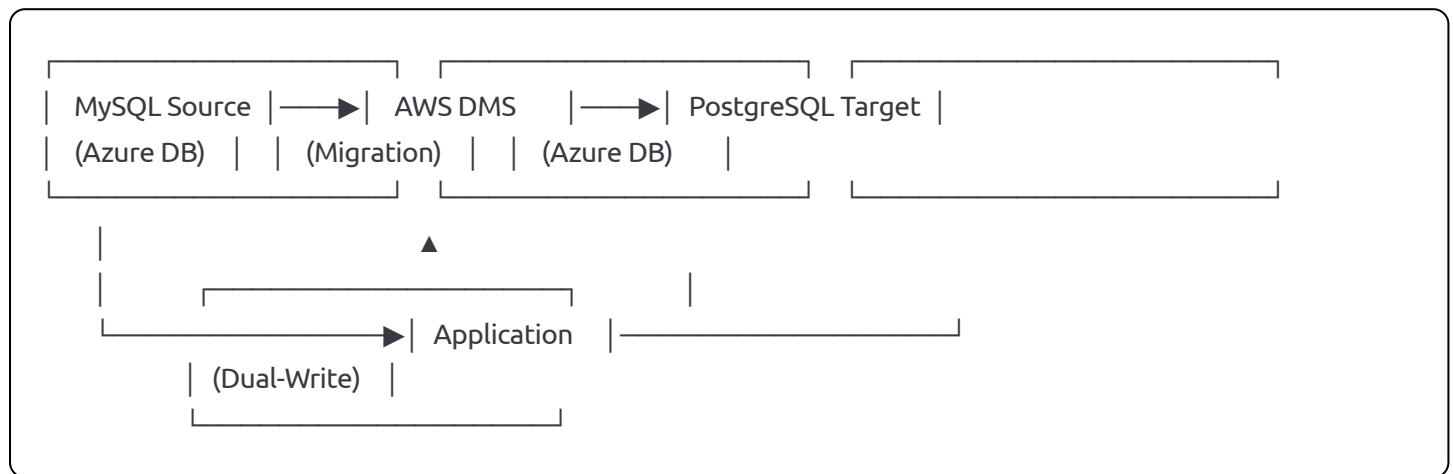


# Zero-Downtime MySQL to PostgreSQL Migration on Azure

## Executive Summary

This guide provides a proven, step-by-step approach for migrating from MySQL to PostgreSQL on Azure with zero downtime using a hybrid strategy combining AWS DMS (for cross-engine support) with Azure-native services.

## Migration Architecture Overview



## Phase 1: Prerequisites and Environment Preparation

### 1.1 Azure Infrastructure Setup

```
bash
```

```
# Create resource group for migration
az group create --name mysql-to-pg-migration --location eastus
```

```
# Create Azure Database for PostgreSQL
```

```
az postgres flexible-server create \
--resource-group mysql-to-pg-migration \
--name myapp-postgresql \
--admin-user pgadmin \
--admin-password 'SecurePassword123!' \
--sku-name Standard_D4s_v3 \
--tier GeneralPurpose \
--storage-size 512 \
--version 14
```

```
# Enable logical replication on source MySQL
```

```
az mysql flexible-server parameter set \
--resource-group mysql-to-pg-migration \
--server-name myapp-mysql \
--name binlog_format \
--value ROW
```

```
az mysql flexible-server parameter set \
--resource-group mysql-to-pg-migration \
--server-name myapp-mysql \
--name binlog_row_image \
--value FULL
```

## 1.2 Network Configuration

```
bash
```

```

# Create VNet for secure communication
az network vnet create \
--resource-group mysql-to-pg-migration \
--name migration-vnet \
--address-prefix 10.0.0.0/16 \
--subnet-name database-subnet \
--subnet-prefix 10.0.1.0/24

# Configure private endpoints
az network private-endpoint create \
--resource-group mysql-to-pg-migration \
--name mysql-private-endpoint \
--vnet-name migration-vnet \
--subnet database-subnet \
--private-connection-resource-id $MYSQL_RESOURCE_ID \
--connection-name mysql-connection \
--group-id mysqlServer

```

## 1.3 MySQL Source Preparation

sql

```

-- Create replication user with necessary privileges
CREATE USER 'replication_user'@'%' IDENTIFIED BY 'ReplicationPass123!';
GRANT REPLICATION SLAVE ON *.* TO 'replication_user'@'%';
GRANT SELECT ON *.* TO 'replication_user'@'%';
GRANT SHOW VIEW ON *.* TO 'replication_user'@'%';
GRANT TRIGGER ON *.* TO 'replication_user'@'%';
FLUSH PRIVILEGES;

-- Enable binary logging (if not already enabled)
-- This requires server restart in Azure MySQL
SET GLOBAL log_bin = ON;
SET GLOBAL binlog_format = 'ROW';
SET GLOBAL binlog_row_image = 'FULL';

-- Verify binary log configuration
SHOW VARIABLES LIKE 'log_bin';
SHOW VARIABLES LIKE 'binlog_format';
SHOW VARIABLES LIKE 'binlog_row_image';

```

## Phase 2: Schema Migration and Conversion

### 2.1 Schema Analysis and Export

bash

```
#!/bin/bash
# schema_export.sh

# Export MySQL schema without data
mysqldump --no-data \
    --routines \
    --triggers \
    --events \
    --single-transaction \
    --host=myapp-mysql.mysql.database.azure.com \
    --user=mysqladmin \
    --password \
    myapp_database > mysql_schema.sql

# Export data for validation
mysqldump --no-create-info \
    --single-transaction \
    --where="1 limit 1000" \
    --host=myapp-mysql.mysql.database.azure.com \
    --user=mysqladmin \
    --password \
    myapp_database > mysql_sample_data.sql
```

## 2.2 Schema Conversion Script

python

```
#!/usr/bin/env python3
# mysql_to_postgresql_converter.py

import re
import sys

def convert_mysql_to_postgresql(mysql_schema):
    """Convert MySQL schema to PostgreSQL compatible schema"""

    conversions = {
        # Data type conversions
        r'\bTINYINT\b': 'BOOLEAN',
        r'\bTINYINT\b': 'SMALLINT',
        r'\bBIGINT\b': 'BIGINT',
        r'\bINT\b': 'UNSIGNED BIGINT',
        r'\bINT\b': 'UNSIGNED INTEGER',
        r'\bINT\b': 'TIMESTAMP',
        r'\bTEXT\b': 'TEXT',
        r'\bLONGTEXT\b': 'TEXT',
        r'\bMEDIUMTEXT\b': 'TEXT',
        r'\bVARCHAR\b': 'VARCHAR',
        r'\bVARCHAR\b': 'CHARACTER VARYING',
        r'\bVARCHAR\b': 'CHARACTER VARYING(1)'

        # AUTO_INCREMENT conversion
        r'\bAUTO_INCREMENT\b': 'GENERATED ALWAYS AS IDENTITY',
    }

    # Engine and charset removal
    r'\s*ENGINE\s*=\s*\w+',
    r'\s*DEFAULT\s*=\s*\w+',
    r'\s*CHARSET\s*=\s*\w+',
    r'\s*COLLATE\s*=\s*\w+',

    # Backtick removal
    r'`[^`]+`': r'"\\1"',

    # Index syntax
    r'\bKEY\s+`?(\\w+)\`?\s*\(([^\`]+)\)`?': r'CREATE INDEX \\1 ON table_name (\\2);',
    r'\bUNIQUE\s+KEY\s+`?(\\w+)\`?\s*\(([^\`]+)\)`?': r'CREATE UNIQUE INDEX \\1 ON table_name (\\2);',
}

postgresql_schema = mysql_schema

for pattern, replacement in conversions.items():
    postgresql_schema = re.sub(pattern, replacement, postgresql_schema, flags=re.IGNORECASE)

# Handle ENUM types
enum_pattern = r"ENUM\s*\(([^\`]+)\)`"
def enum_replacement(match):
```

```

values = match.group(1)
return f"VARCHAR(255) CHECK (column_name IN ({values}))"

postgresql_schema = re.sub(enum_pattern, enum_replacement, postgresql_schema, flags=re.IGNORECASE)

return postgresql_schema

def main():
    if len(sys.argv) != 3:
        print("Usage: python3 mysql_to_postgresql_converter.py input.sql output.sql")
        sys.exit(1)

    input_file = sys.argv[1]
    output_file = sys.argv[2]

    with open(input_file, 'r') as f:
        mysql_schema = f.read()

    postgresql_schema = convert_mysql_to_postgresql(mysql_schema)

    with open(output_file, 'w') as f:
        f.write(postgresql_schema)

    print(f"Schema conversion completed. Output saved to {output_file}")

if __name__ == "__main__":
    main()

```

## 2.3 Execute Schema Creation

```

bash

# Convert schema
python3 mysql_to_postgresql_converter.py mysql_schema.sql postgresql_schema.sql

# Create schema in PostgreSQL
PGPASSWORD='SecurePassword123!' psql \
-h myapp-postgresql.postgres.database.azure.com \
-U pgadmin \
-d myapp_database \
-f postgresql_schema.sql

```

## Phase 3: AWS DMS Setup for Cross-Engine Migration

### 3.1 AWS DMS Infrastructure

bash

```
# Create AWS DMS replication instance
aws dms create-replication-instance \
--replication-instance-identifier mysql-to-postgresql-migration \
--replication-instance-class dms.t3.medium \
--allocated-storage 20 \
--vpc-security-group-ids sg-xxxxxxxxxx \
--replication-subnet-group-identifier default-vpc-xxxxxxxx \
--publicly-accessible \
--multi-az false
```

```
# Create source endpoint (MySQL on Azure)
```

```
aws dms create-endpoint \
--endpoint-identifier mysql-source-azure \
--endpoint-type source \
--engine-name mysql \
--server-name myapp-mysql.mysql.database.azure.com \
--port 3306 \
--username mysqladmin \
--password 'MySQLPassword123!' \
--database-name myapp_database
```

```
# Create target endpoint (PostgreSQL on Azure)
```

```
aws dms create-endpoint \
--endpoint-identifier postgresql-target-azure \
--endpoint-type target \
--engine-name postgres \
--server-name myapp-postgresql.postgres.database.azure.com \
--port 5432 \
--username pgadmin \
--password 'SecurePassword123!' \
--database-name myapp_database
```

## 3.2 DMS Task Configuration

json

```
{
  "rules": [
    {
      "rule-type": "selection",
      "rule-id": "1",
      "rule-name": "1",
      "object-locator": {
        "schema-name": "myapp_database",
        "table-name": "%"
      },
      "rule-action": "include",
      "filters": []
    },
    {
      "rule-type": "transformation",
      "rule-id": "2",
      "rule-name": "2",
      "rule-target": "column",
      "object-locator": {
        "schema-name": "myapp_database",
        "table-name": "%",
        "column-name": "%"
      },
      "rule-action": "convert-lowercase"
    }
  ]
}
```

bash

```
# Create migration task
aws dms create-replication-task \
--replication-task-identifier mysql-to-postgresql-continuous \
--source-endpoint-arn arn:aws:dms:us-east-1:account:endpoint:mysql-source-azure \
--target-endpoint-arn arn:aws:dms:us-east-1:account:endpoint:postgresql-target-azure \
--replication-instance-arn arn:aws:dms:us-east-1:account:rep:mysql-to-postgresql-migration \
--migration-type full-load-and-cdc \
--table-mappings file://table-mappings.json \
--replication-task-settings file://task-settings.json
```

## Phase 4: Alternative CDC Setup (Debezium + Kafka)

### 4.1 Kafka and Debezium Setup

yaml

```

# docker-compose.yml for Kafka + Debezium
version: '3.8'
services:
zookeeper:
  image: confluentinc/cp-zookeeper:latest
  environment:
    ZOOKEEPER_CLIENT_PORT: 2181
    ZOOKEEPER_TICK_TIME: 2000

kafka:
  image: confluentinc/cp-kafka:latest
  depends_on:
    - zookeeper
  ports:
    - "9092:9092"
  environment:
    KAFKA_BROKER_ID: 1
    KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
    KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://localhost:9092
    KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1

debezium:
  image: debezium/connect:latest
  depends_on:
    - kafka
  ports:
    - "8083:8083"
  environment:
    BOOTSTRAP_SERVERS: kafka:9092
    GROUP_ID: 1
    CONFIG_STORAGE_TOPIC: debezium_configs
    OFFSET_STORAGE_TOPIC: debezium_offsets
    STATUS_STORAGE_TOPIC: debezium_statuses

```

## 4.2 Debezium MySQL Connector Configuration

json

```
{  
  "name": "mysql-postgres-connector",  
  "config": {  
    "connector.class": "io.debezium.connector.mysql.MySqlConnector",  
    "database.hostname": "myapp-mysql.mysql.database.azure.com",  
    "database.port": "3306",  
    "database.user": "replication_user",  
    "database.password": "ReplicationPass123!",  
    "database.server.id": "12345",  
    "database.server.name": "azure-mysql-server",  
    "database.include.list": "myapp_database",  
    "table.include.list": "myapp_database.users,myapp_database.orders,myapp_database.products",  
    "database.history.kafka.bootstrap.servers": "kafka:9092",  
    "database.history.kafka.topic": "dbhistory.azure-mysql-server",  
    "include.schema.changes": "true",  
    "transforms": "unwrap",  
    "transforms.unwrap.type": "io.debezium.transforms.ExtractNewRecordState",  
    "transforms.unwrap.drop.tombstones": "false",  
    "key.converter": "org.apache.kafka.connect.json.JsonConverter",  
    "value.converter": "org.apache.kafka.connect.json.JsonConverter",  
    "key.converter.schemas.enable": "false",  
    "value.converter.schemas.enable": "false"  
  }  
}
```

## 4.3 Kafka to PostgreSQL Sync Service

python

```
#!/usr/bin/env python3
# kafka_to_postgres_sync.py

import json
import asyncio
import asyncpg
from kafka import KafkaConsumer
from typing import Dict, Any
import logging

logging.basicConfig(level=logging.INFO)
logger = logging.getLogger(__name__)

class KafkaPostgreSync:
    def __init__(self, kafka_config: Dict, postgres_config: Dict):
        self.kafka_config = kafka_config
        self.postgres_config = postgres_config
        self.pool = None

    async def init_postgres_pool(self):
        """Initialize PostgreSQL connection pool"""
        self.pool = await asyncpg.create_pool(
            host=self.postgres_config['host'],
            port=self.postgres_config['port'],
            user=self.postgres_config['user'],
            password=self.postgres_config['password'],
            database=self.postgres_config['database'],
            min_size=5,
            max_size=20
        )

    def create_kafka_consumer(self) -> KafkaConsumer:
        """Create Kafka consumer"""
        return KafkaConsumer(
            'azure-mysql-server.myapp_database.users',
            'azure-mysql-server.myapp_database.orders',
            'azure-mysql-server.myapp_database.products',
            bootstrap_servers=self.kafka_config['bootstrap_servers'],
            value_deserializer=lambda m: json.loads(m.decode('utf-8')),
            key_deserializer=lambda m: json.loads(m.decode('utf-8')) if m else None,
            auto_offset_reset='earliest',
            enable_auto_commit=True,
            group_id='postgres-sync-group'
        )

    async def process_message(self, message):
```

```

"""Process Kafka message and sync to PostgreSQL"""
try:
    # Extract table name from topic
    topic_parts = message.topic.split('.')
    table_name = topic_parts[-1]

    # Get message payload
    payload = message.value
    operation = payload.get('op')

    if operation == 'c': # Create
        await self.handle_insert(table_name, payload['after'])
    elif operation == 'u': # Update
        await self.handle_update(table_name, payload['before'], payload['after'])
    elif operation == 'd': # Delete
        await self.handle_delete(table_name, payload['before'])
    elif operation == 'r': # Read (initial snapshot)
        await self.handle_insert(table_name, payload['after'])

except Exception as e:
    logger.error(f"Error processing message: {e}")

async def handle_insert(self, table_name: str, data: Dict[str, Any]):
    """Handle INSERT operation"""
    async with self.pool.acquire() as conn:
        columns = list(data.keys())
        values = list(data.values())
        placeholders = [f"${i+1}" for i in range(len(values))]

        query = f"""
            INSERT INTO {table_name} ({', '.join(columns)})
            VALUES ({', '.join(placeholders)})
            ON CONFLICT DO NOTHING
        """
        await conn.execute(query, *values)
        logger.info(f"Inserted record into {table_name}")

async def handle_update(self, table_name: str, before: Dict[str, Any], after: Dict[str, Any]):
    """Handle UPDATE operation"""
    async with self.pool.acquire() as conn:
        # Assume 'id' is primary key
        where_clause = "id = $1"
        set_clauses = []
        values = [after['id']]

        param_counter = 2

```

```
for key, value in after.items():
    if key != 'id':
        set_clauses.append(f'{key} = ${param_counter}')
        values.append(value)
        param_counter += 1

query = f"""
UPDATE {table_name}
SET {', '.join(set_clauses)}
WHERE {where_clause}
"""

await conn.execute(query, *values)
logger.info(f"Updated record in {table_name}")

async def handle_delete(self, table_name: str, data: Dict[str, Any]):
    """Handle DELETE operation"""
    async with self.pool.acquire() as conn:
        query = f"DELETE FROM {table_name} WHERE id = $1"
        await conn.execute(query, data['id'])
        logger.info(f"Deleted record from {table_name}")

async def start_sync(self):
    """Start the synchronization process"""
    await self.init_postgres_pool()
    consumer = self.create_kafka_consumer()

    logger.info("Starting Kafka to PostgreSQL sync... ")

    try:
        for message in consumer:
            await self.process_message(message)
    except KeyboardInterrupt:
        logger.info("Stopping sync...")
    finally:
        consumer.close()
        await self.pool.close()

# Configuration
kafka_config = {
    'bootstrap_servers': ['localhost:9092']
}

postgres_config = {
    'host': 'myapp-postgresql.postgres.database.azure.com',
    'port': 5432,
    'user': 'pgladmin',
```

```
'password': 'SecurePassword123!',  
'database': 'myapp_database'  
}  
  
# Run the sync service  
async def main():  
    sync_service = KafkaPostgreSQLSync(kafka_config, postgres_config)  
    await sync_service.start_sync()  
  
if __name__ == "__main__":  
    asyncio.run(main())
```

## Phase 5: Application Dual-Write Strategy

### 5.1 Database Abstraction Layer

```
python
```

```
#!/usr/bin/env python3
# database_abstraction.py

import asyncio
import asyncpg
import aiomysql
from typing import Optional, Dict, Any, List
import logging
from contextlib import asynccontextmanager

logger = logging.getLogger(__name__)

class DatabaseManager:
    def __init__(self, mysql_config: Dict, postgres_config: Dict):
        self.mysql_config = mysql_config
        self.postgres_config = postgres_config
        self.mysql_pool = None
        self.postgres_pool = None
        self.dual_write_enabled = False
        self.primary_db = 'mysql' # Start with MySQL as primary

    async def init_connections(self):
        """Initialize database connections"""
        # MySQL connection
        self.mysql_pool = await aiomysql.create_pool(
            host=self.mysql_config['host'],
            port=self.mysql_config['port'],
            user=self.mysql_config['user'],
            password=self.mysql_config['password'],
            db=self.mysql_config['database'],
            minsize=5,
            maxsize=20,
            autocommit=False
        )

        # PostgreSQL connection
        self.postgres_pool = await asyncpg.create_pool(
            host=self.postgres_config['host'],
            port=self.postgres_config['port'],
            user=self.postgres_config['user'],
            password=self.postgres_config['password'],
            database=self.postgres_config['database'],
            min_size=5,
            max_size=20
        )
```

```
logger.info("Database connections initialized")

def enable_dual_write(self):
    """Enable dual-write mode"""
    self.dual_write_enabled = True
    logger.info("Dual-write mode enabled")

def switch_primary(self, primary: str):
    """Switch primary database"""
    if primary in ['mysql', 'postgres']:
        self.primary_db = primary
        logger.info(f"Primary database switched to {primary}")
    else:
        raise ValueError("Primary must be 'mysql' or 'postgres')

async def execute_read(self, query: str, params: Optional[List] = None) -> List[Dict[str, Any]]:
    """Execute read query on primary database"""
    if self.primary_db == 'mysql':
        return await self._execute_mysql_read(query, params)
    else:
        return await self._execute_postgres_read(query, params)

async def execute_write(self, query: str, params: Optional[List] = None) -> bool:
    """Execute write query with dual-write support"""
    success = True

    if self.primary_db == 'mysql':
        success &= await self._execute_mysql_write(query, params)
        if self.dual_write_enabled:
            # Convert MySQL query to PostgreSQL
            pg_query = self._convert_mysql_to_postgres(query)
            success &= await self._execute_postgres_write(pg_query, params)
    else:
        success &= await self._execute_postgres_write(query, params)
        if self.dual_write_enabled:
            # Convert PostgreSQL query to MySQL
            mysql_query = self._convert_postgres_to_mysql(query)
            success &= await self._execute_mysql_write(mysql_query, params)

    return success

async def _execute_mysql_read(self, query: str, params: Optional[List] = None) -> List[Dict[str, Any]]:
    """Execute MySQL read query"""
    async with self.mysql_pool.acquire() as conn:
        async with conn.cursor(aiomysql.DictCursor) as cursor:
            await cursor.execute(query, params or [])
            results = await cursor.fetchall()
```

```
        return list(results)

async def _execute_mysql_write(self, query: str, params: Optional[List] = None) -> bool:
    """Execute MySQL write query"""
    try:
        async with self.mysql_pool.acquire() as conn:
            async with conn.cursor() as cursor:
                await cursor.execute(query, params or [])
                await conn.commit()
        return True
    except Exception as e:
        logger.error(f"MySQL write error: {e}")
        return False

async def _execute_postgres_read(self, query: str, params: Optional[List] = None) -> List[Dict[str, Any]]:
    """Execute PostgreSQL read query"""
    async with self.postgres_pool.acquire() as conn:
        rows = await conn.fetch(query, *(params or []))
    return [dict(row) for row in rows]

async def _execute_postgres_write(self, query: str, params: Optional[List] = None) -> bool:
    """Execute PostgreSQL write query"""
    try:
        async with self.postgres_pool.acquire() as conn:
            async with conn.transaction():
                await conn.execute(query, *(params or []))
        return True
    except Exception as e:
        logger.error(f"PostgreSQL write error: {e}")
        return False

def _convert_mysql_to_postgres(self, query: str) -> str:
    """Convert MySQL query syntax to PostgreSQL"""
    # Basic conversions
    conversions = {
        'LIMIT %s, %s': 'LIMIT %s OFFSET %s',
        '%': '',
        'AUTO_INCREMENT': 'GENERATED ALWAYS AS IDENTITY'
    }

    converted_query = query
    for mysql_syntax, postgres_syntax in conversions.items():
        converted_query = converted_query.replace(mysql_syntax, postgres_syntax)

    return converted_query

def _convert_postgres_to_mysql(self, query: str) -> str:
```

```

"""Convert PostgreSQL query syntax to MySQL"""
# Basic conversions
conversions = {
    'LIMIT %s OFFSET %s': 'LIMIT %s, %s',
    '": ":"',
    'GENERATED ALWAYS AS IDENTITY': 'AUTO_INCREMENT'
}

converted_query = query
for postgres_syntax, mysql_syntax in conversions.items():
    converted_query = converted_query.replace(postgres_syntax, mysql_syntax)

return converted_query

async def validate_data_consistency(self, table_name: str) -> Dict[str, Any]:
    """Validate data consistency between MySQL and PostgreSQL"""

    # Count rows
    mysql_count_query = f"SELECT COUNT(*) as count FROM `'{table_name}'`"
    postgres_count_query = f"SELECT COUNT(*) as count FROM \"'{table_name}'\""

    mysql_count = await self._execute_mysql_read(mysql_count_query)
    postgres_count = await self._execute_postgres_read(postgres_count_query)

    # Checksum comparison (simplified)
    mysql_checksum_query = f"""
        SELECT BIT_XOR(CRC32(CONCAT_WS('|',
            COALESCE(id, ''),
            COALESCE(name, ''),
            COALESCE(email, '')
        ))) as checksum
        FROM `'{table_name}'`"""
    """

    postgres_checksum_query = f"""
        SELECT BIT_XOR(
            ('x' || substr(md5(COALESCE(id::text, "") || '|' ||
                COALESCE(name, '') || '|' ||
                COALESCE(email, '')), 1, 8)::bit(32)::int
        ) as checksum
        FROM \"'{table_name}'\""""
    """

try:
    mysql_checksum = await self._execute_mysql_read(mysql_checksum_query)
    postgres_checksum = await self._execute_postgres_read(postgres_checksum_query)
except Exception as e:
    logger.warning(f"Checksum comparison failed: {e}")

```

```
mysql_checksum = [{'checksum': 0}]
postgres_checksum = [{'checksum': 0}]

return {
    'table_name': table_name,
    'mysql_count': mysql_count[0]['count'],
    'postgres_count': postgres_count[0]['count'],
    'count_match': mysql_count[0]['count'] == postgres_count[0]['count'],
    'mysql_checksum': mysql_checksum[0]['checksum'],
    'postgres_checksum': postgres_checksum[0]['checksum'],
    'checksum_match': mysql_checksum[0]['checksum'] == postgres_checksum[0]['checksum']
}

async def close_connections(self):
    """Close database connections"""
    if self.mysql_pool:
        self.mysql_pool.close()
        await self.mysql_pool.wait_closed()

    if self.postgres_pool:
        await self.postgres_pool.close()

    logger.info("Database connections closed")
```

## 5.2 Application Integration Example

python

```
#!/usr/bin/env python3
# application_service.py

import asyncio
from database_abstraction import DatabaseManager
from typing import Optional, Dict, Any, List

class UserService:
    def __init__(self, db_manager: DatabaseManager):
        self.db = db_manager

    async def create_user(self, user_data: Dict[str, Any]) -> Optional[int]:
        """Create a new user"""
        query = """
            INSERT INTO users (name, email, created_at)
            VALUES (%s, %s, NOW())
        """

        success = await self.db.execute_write(
            query,
            [user_data['name'], user_data['email']]
        )

        if success:
            # Get the ID (this is simplified)
            result = await self.db.execute_read(
                "SELECT LAST_INSERT_ID() as id" if self.db.primary_db == 'mysql'
                else "SELECT lastval() as id"
            )
            return result[0]['id']

        return None

    async def get_user(self, user_id: int) -> Optional[Dict[str, Any]]:
        """Get user by ID"""
        query = "SELECT * FROM users WHERE id = %s"
        results = await self.db.execute_read(query, [user_id])
        return results[0] if results else None

    async def update_user(self, user_id: int, user_data: Dict[str, Any]) -> bool:
        """Update user"""
        query = """
            UPDATE users
            SET name = %s, email = %s, updated_at = NOW()
            WHERE id = %s
        """

        success = await self.db.execute_write(
            query,
            [user_data['name'], user_data['email'], user_id]
        )

        return success
```

```
return await self.db.execute_write(
    query,
    [user_data['name'], user_data['email'], user_id]
)

async def delete_user(self, user_id: int) -> bool:
    """Delete user"""
    query = "DELETE FROM users WHERE id = %s"
    return await self.db.execute_write(query, [user_id])

# Usage example
async def main():
    # Database configurations
    mysql_config = {
        'host': 'myapp-mysql.mysql.database.azure.com',
        'port': 3306,
        'user': 'mysqladmin',
        'password': 'MySQLPassword123!',
        'database': 'myapp_database'
    }

    postgres_config = {
        'host': 'myapp-postgresql.postgres.database.azure.com',
        'port': 5432,
        'user': 'pgadmin',
        'password': 'SecurePassword123!',
        'database': 'myapp_database'
    }

    # Initialize database manager
    db_manager = DatabaseManager(mysql_config, postgres_config)
    await db_manager.init_connections()

    # Enable dual-write mode during migration
    db_manager.enable_dual_write()

    # Initialize services
    user_service = UserService(db_manager)

    # Test operations
    user_id = await user_service.create_user({
        'name': 'John Doe',
        'email': 'john@example.com'
    })

    print(f"Created user with ID: {user_id}")
```

```
# Validate consistency
consistency_report = await db_manager.validate_data_consistency('users')
print(f"Consistency report: {consistency_report}")

# After migration is complete, switch to PostgreSQL
db_manager.switch_primary('postgres')

# Continue operations on PostgreSQL
user = await user_service.get_user(user_id)
print(f"Retrieved user: {user}")

await db_manager.close_connections()

if __name__ == "__main__":
    asyncio.run(main())
```

## Phase 6: Testing and Validation

### 6.1 Data Validation Scripts

bash

```
#!/bin/bash
# validation_suite.sh

echo "Starting comprehensive data validation..."

# Row count validation
python3 << EOF
import asyncio
import asyncpg
import aiomysql

async def validate_row_counts():
    # MySQL connection
    mysql_conn = await aiomysql.connect(
        host='myapp-mysql.mysql.database.azure.com',
        port=3306,
        user='mysqladmin',
        password='MySQLPassword123!',
        db='myapp_database'
    )

    # PostgreSQL connection
    pg_conn = await asyncpg.connect(
        host='myapp-postgresql.postgres.database.azure.com',
        port=5432,
        user='pgladmin',
        password='SecurePassword123!',
        database='myapp_database'
    )

    tables = ['users', 'orders', 'products', 'order_items']

    for table in tables:
        # MySQL count
        async with mysql_conn.cursor() as cursor:
            await cursor.execute(f"SELECT COUNT(*) FROM {table}")
            mysql_count = await cursor.fetchone()

        # PostgreSQL count
        pg_count = await pg_conn.fetchval(f"SELECT COUNT(*) FROM \"{table}\"")

        print(f"Table {table}:")
        print(f"  MySQL: {mysql_count[0]}")
        print(f"  PostgreSQL: {pg_count}")
        print(f"  Match: {mysql_count[0] == pg_count}")
        print()
```

```
mysql_conn.close()
await pg_conn.close()

asyncio.run(validate_row_counts())
EOF

echo "Row count validation completed."

# Schema validation
echo "Validating schema differences..."
python3 << EOF
import asyncio
import asyncpg
import aiomysql

async def validate_schemas():
    # MySQL schema info
    mysql_conn = await aiomysql.connect(
        host='myapp-mysql.mysql.database.azure.com',
        port=3306,
        user='mysqladmin',
        password='MySQLPassword123!',
        db='myapp_database'
    )

    # PostgreSQL schema info
    pg_conn = await asyncpg.connect(
        host='myapp-postgresql.postgres.database.azure.com',
        port=5432,
        user='pgadmin',
        password='SecurePassword123!',
        database='myapp_database'
    )

    # Get MySQL table structure
    async with mysql_conn.cursor(aiomysql.DictCursor) as cursor:
        await cursor.execute("""
            SELECT TABLE_NAME, COLUMN_NAME, DATA_TYPE, IS_NULLABLE, COLUMN_DEFAULT
            FROM INFORMATION_SCHEMA.COLUMNS
            WHERE TABLE_SCHEMA = 'myapp_database'
            ORDER BY TABLE_NAME, ORDINAL_POSITION
        """)
        mysql_schema = await cursor.fetchall()

    # Get PostgreSQL table structure
    pg_schema = await pg_conn.fetch("""

```

```
SELECT table_name, column_name, data_type, is_nullable, column_default
FROM information_schema.columns
WHERE table_schema = 'public'
ORDER BY table_name, ordinal_position
""")  
  
print("Schema comparison:")  
print("MySQL tables:", {row['TABLE_NAME'] for row in mysql_schema})  
print("PostgreSQL tables:", {row['table_name'] for row in pg_schema})  
  
mysql_conn.close()  
await pg_conn.close()  
  
asyncio.run(validate_schemas())  
EOF  
  
echo "Schema validation completed."
```

## 6.2 Performance Testing

```
python
```



```

start_time = time.time()
db_manager.primary_db = 'mysql'
await db_manager.execute_read(query)
mysql_times.append(time.time() - start_time)

# Test PostgreSQL
postgres_times = []
for _ in range(10):
    start_time = time.time()
    db_manager.primary_db = 'postgres'
    await db_manager.execute_read(query)
    postgres_times.append(time.time() - start_time)

results[query] = {
    'mysql_avg': statistics.mean(mysql_times),
    'mysql_median': statistics.median(mysql_times),
    'postgres_avg': statistics.mean(postgres_times),
    'postgres_median': statistics.median(postgres_times)
}

# Print results
print("\nPerformance Test Results:")
print("=" * 80)

for query, metrics in results.items():
    print(f"\nQuery: {query[:50]}...")
    print(f"MySQL - Avg: {metrics['mysql_avg']:.4f}s, Median: {metrics['mysql_median']:.4f}s")
    print(f"PostgreSQL - Avg: {metrics['postgres_avg']:.4f}s, Median: {metrics['postgres_median']:.4f}s")

improvement = ((metrics['mysql_avg'] - metrics['postgres_avg']) / metrics['mysql_avg']) * 100
print(f"Performance change: {improvement:+.2f}%")

await db_manager.close_connections()

if __name__ == "__main__":
    asyncio.run(performance_test())

```

## Phase 7: Cutover Process

### 7.1 Pre-Cutover Checklist

bash

```
#!/bin/bash
# pre_cutover_checklist.sh

echo "Pre-Cutover Validation Checklist"
echo "====="

# 1. Verify replication lag
echo "1. Checking replication lag..."
aws dms describe-replication-tasks \
--filters Name=replication-task-id,Values=mysql-to-postgresql-continuous \
--query 'ReplicationTasks[0].ReplicationTaskStats'

# 2. Validate data consistency
echo "2. Running data consistency checks..."
python3 validation_suite.sh

# 3. Test PostgreSQL performance
echo "3. Running performance tests..."
python3 performance_test.py

# 4. Verify backup status
echo "4. Checking backup status..."
az postgres flexible-server backup list \
--resource-group mysql-to-pg-migration \
--server-name myapp-postgresql

# 5. Test application connectivity
echo "5. Testing application connectivity..."
python3 << EOF
import asyncio
from database_abstraction import DatabaseManager

async def test_connectivity():
    mysql_config = {
        'host': 'myapp-mysql.mysql.database.azure.com',
        'port': 3306,
        'user': 'mysqladmin',
        'password': 'MySQLPassword123!',
        'database': 'myapp_database'
    }

    postgres_config = {
        'host': 'myapp-postgresql.postgres.database.azure.com',
        'port': 5432,
        'user': 'pgadmin',
        'password': 'SecurePassword123!',
```

```
'database': 'myapp_database'
}

db_manager = DatabaseManager(mysql_config, postgres_config)
await db_manager.init_connections()

# Test both connections
try:
    mysql_result = await db_manager._execute_mysql_read("SELECT 1 as test")
    print(f"MySQL connection: ✓ {mysql_result}")
except Exception as e:
    print(f"MySQL connection: ✗ {e}")

try:
    postgres_result = await db_manager._execute_postgres_read("SELECT 1 as test")
    print(f"PostgreSQL connection: ✓ {postgres_result}")
except Exception as e:
    print(f"PostgreSQL connection: ✗ {e}")

await db_manager.close_connections()

asyncio.run(test_connectivity())
EOF

echo "Pre-cutover validation completed."
```

## 7.2 Cutover Script

bash

```
#!/bin/bash
# cutover_script.sh

set -e # Exit on any error

echo "Starting MySQL to PostgreSQL Cutover Process"
echo "====="

# Configuration
MAINTENANCE_START=$(date -u +"%Y-%m-%d %H:%M:%S UTC")
ROLLBACK_TIMEOUT=3600 # 1 hour rollback window

echo "Cutover started at: $MAINTENANCE_START"

# 1. Enable maintenance mode
echo "1. Enabling maintenance mode..."
# Update load balancer to show maintenance page
az network application-gateway http-settings update \
--resource-group mysql-to-pg-migration \
--gateway-name myapp-gateway \
--name maintenance-settings \
--path /maintenance.html

# 2. Stop application writes
echo "2. Stopping application writes..."
# Scale down application instances or enable read-only mode
kubectl scale deployment myapp --replicas=0

# Wait for active connections to drain
echo "Waiting 30 seconds for connections to drain..."
sleep 30

# 3. Verify replication is caught up
echo "3. Verifying replication lag..."
python3 << EOF
import boto3
import time

dms_client = boto3.client('dms')

while True:
    response = dms_client.describe_replication_tasks(
        Filters=[

    {
        'Name': 'replication-task-id',
        'Values': ['mysql-to-postgresql-continuous']
    }
]
)
```

```
        }
    ]
}

task = response['ReplicationTasks'][0]
stats = task['ReplicationTaskStats']

lag_seconds = stats.get('Elapsed Time Millis', 0) / 1000
print(f"Current replication lag: {lag_seconds} seconds")

if lag_seconds < 5: # Less than 5 seconds lag
    print("Replication is caught up!")
    break

time.sleep(10)
EOF

# 4. Final data validation
echo "4. Running final data validation..."
python3 validation_suite.sh

# 5. Stop replication
echo "5. Stopping replication..."
aws dms stop-replication-task \
--replication-task-arn arn:aws:dms:us-east-1:account:task:mysql-to-postgresql-continuous

# 6. Switch application to PostgreSQL
echo "6. Switching application to PostgreSQL..."

# Update application configuration
kubectl create configmap myapp-config \
--from-literal=DATABASE_URL="postgresql://pgadmin:SecurePassword123!@myapp-postgresql.postgres.database.amazonaws.com:5432/myapp" \
--dry-run=client -o yaml | kubectl apply -f -

# Update deployment
kubectl set env deployment/myapp PRIMARY_DB=postgres
kubectl set env deployment/myapp DUAL_WRITE_ENABLED=false

# Scale up application
kubectl scale deployment myapp --replicas=3

# 7. Test application functionality
echo "7. Testing application functionality..."
sleep 60 # Wait for pods to be ready

# Health check
for i in {1..10}; do
```

```

HTTP_STATUS=$(curl -s -o /dev/null -w "%{http_code}" https://myapp.example.com/health)
if [ "$HTTP_STATUS" = "200" ]; then
    echo "Application health check passed"
    break
else
    echo "Health check failed (attempt $i/10): HTTP $HTTP_STATUS"
    if [ $i -eq 10 ]; then
        echo "Health checks failed, initiating rollback..."
        bash rollback_script.sh
        exit 1
    fi
    sleep 10
fi
done

# 8. Disable maintenance mode
echo "8. Disabling maintenance mode..."
az network application-gateway http-settings update \
--resource-group mysql-to-pg-migration \
--gateway-name myapp-gateway \
--name default-settings \
--path /

# 9. Monitor for issues
echo "9. Monitoring application..."
echo "Cutover completed successfully at: $(date -u +"%Y-%m-%d %H:%M:%S UTC")"
echo "Monitoring for 30 minutes before declaring success..."

# Monitor application metrics for 30 minutes
for i in {1..30}; do
    echo "Monitoring minute $i/30..."

# Check error rates, response times, etc.
ERROR_COUNT=$(kubectl logs deployment/myapp --since=1m | grep -c "ERROR" || echo "0")
if [ "$ERROR_COUNT" -gt 10 ]; then
    echo "High error rate detected: $ERROR_COUNT errors in the last minute"
    echo "Consider investigating or rolling back"
fi

sleep 60
done

echo "Cutover monitoring completed successfully!"
echo "Migration from MySQL to PostgreSQL is complete."

```

```
# 10. Clean up old resources (optional, run later)
echo "Remember to clean up MySQL resources after confirming stability"
```

## 7.3 Rollback Script

bash

```
#!/bin/bash
# rollback_script.sh

set -e

echo "EMERGENCY ROLLBACK: Switching back to MySQL"
echo "====="

# 1. Enable maintenance mode
echo "1. Enabling maintenance mode..."
az network application-gateway http-settings update \
--resource-group mysql-to-pg-migration \
--gateway-name myapp-gateway \
--name maintenance-settings \
--path /maintenance.html

# 2. Scale down application
echo "2. Scaling down application..."
kubectl scale deployment myapp --replicas=0

# 3. Switch back to MySQL configuration
echo "3. Switching back to MySQL..."
kubectl create configmap myapp-config \
--from-literal=DATABASE_URL="mysql://mysqladmin:MySQLPassword123!@myapp-mysql.mysql.database.azure.com" \
--dry-run=client -o yaml | kubectl apply -f -

kubectl set env deployment/myapp PRIMARY_DB=mysql
kubectl set env deployment/myapp DUAL_WRITE_ENABLED=true

# 4. Scale up application
echo "4. Scaling up application..."
kubectl scale deployment myapp --replicas=3

# 5. Wait for readiness
echo "5. Waiting for application readiness..."
kubectl wait --for=condition=available --timeout=300s deployment/myapp

# 6. Test functionality
echo "6. Testing application..."
for i in {1..5}; do
    HTTP_STATUS=$(curl -s -o /dev/null -w "%{http_code}" https://myapp.example.com/health)
    if [ "$HTTP_STATUS" = "200" ]; then
        echo "Application health check passed"
        break
    fi
    sleep 10
done
```

done

```
# 7. Disable maintenance mode
echo "7. Disabling maintenance mode..."
az network application-gateway http-settings update \
--resource-group mysql-to-pg-migration \
--gateway-name myapp-gateway \
--name default-settings \
--path /
echo "Rollback completed successfully!"
echo "Application is now running on MySQL again."
```

## Phase 8: Post-Migration Monitoring

### 8.1 Monitoring Setup

yaml

```
# monitoring-stack.yaml
apiVersion: v1
kind: ConfigMap
metadata:
  name: prometheus-config
data:
  prometheus.yml: |
    global:
      scrape_interval: 15s
    scrape_configs:
      - job_name: 'postgresql-exporter'
        static_configs:
          - targets: ['postgres-exporter:9187']
      - job_name: 'application'
        static_configs:
          - targets: ['myapp:8080']

    rule_files:
      - "alert_rules.yml"

  alerting:
    alertmanagers:
      - static_configs:
          - targets: ['alertmanager:9093']

  alert_rules.yml: |
    groups:
      - name: database_alerts
        rules:
          - alert: PostgreSQLDown
            expr: up{job="postgresql-exporter"} == 0
            for: 1m
            labels:
              severity: critical
            annotations:
              summary: "PostgreSQL is down"
              description: "PostgreSQL database is not responding"

          - alert: HighErrorRate
            expr: rate(http_requests_total{status=~"5.."}[5m]) > 0.1
            for: 2m
            labels:
              severity: warning
            annotations:
              summary: "High error rate detected"
              description: "Error rate is {{ $value }} errors per second"
```

```
- alert: SlowQueries
  expr: postgresql_stat_activity_max_tx_duration{state="active"} > 300
  for: 5m
  labels:
    severity: warning
  annotations:
    summary: "Long running queries detected"
    description: "Query running for more than 5 minutes"
```

## 8.2 Performance Monitoring Script

python

```
#!/usr/bin/env python3
# post_migration_monitor.py

import asyncio
import asyncpg
import time
import json
import logging
from datetime import datetime, timedelta
from typing import Dict, List, Any

logging.basicConfig(level=logging.INFO)
logger = logging.getLogger(__name__)

class PostgreSQLMonitor:
    def __init__(self, connection_config: Dict[str, str]):
        self.config = connection_config
        self.pool = None

    async def init_connection(self):
        """Initialize PostgreSQL connection pool"""
        self.pool = await asyncpg.create_pool(
            host=self.config['host'],
            port=self.config['port'],
            user=self.config['user'],
            password=self.config['password'],
            database=self.config['database'],
            min_size=2,
            max_size=5
        )

    async def get_database_metrics(self) -> Dict[str, Any]:
        """Get database performance metrics"""
        async with self.pool.acquire() as conn:
            metrics = {}

            # Connection count
            metrics['connections'] = await conn.fetchval("""
                SELECT count(*) FROM pg_stat_activity
                WHERE state = 'active' AND pid != pg_backend_pid()
            """)

            # Database size
            metrics['database_size_mb'] = await conn.fetchval("""
                SELECT pg_size_pretty(pg_database_size(current_database()))
            """)

            return metrics
```

```

# Cache hit ratio
metrics['cache_hit_ratio'] = await conn.fetchval("""
    SELECT round(
        100 * sum(blks_hit) / (sum(blks_hit) + sum(blks_read)), 2
    ) as cache_hit_ratio
    FROM pg_stat_database
    WHERE datname = current_database()
""")

# Long running queries
long_queries = await conn.fetch("""
    SELECT pid, now() - pg_stat_activity.query_start AS duration, query
    FROM pg_stat_activity
    WHERE (now() - pg_stat_activity.query_start) > interval '5 minutes'
    AND state = 'active'
    AND pid != pg_backend_pid()
""")
metrics['long_queries'] = len(long_queries)

# Table statistics
table_stats = await conn.fetch("""
    SELECT schemaname, tablename, n_tup_ins, n_tup_upd, n_tup_del,
        n_tup_hot_upd, n_live_tup, n_dead_tup
    FROM pg_stat_user_tables
    ORDER BY n_live_tup DESC
    LIMIT 10
""")
metrics['top_tables'] = [dict(row) for row in table_stats]

# Index usage
index_usage = await conn.fetch("""
    SELECT schemaname, tablename, indexname, idx_tup_read, idx_tup_fetch
    FROM pg_stat_user_indexes
    WHERE idx_tup_read > 0
    ORDER BY idx_tup_read DESC
    LIMIT 10
""")
metrics['index_usage'] = [dict(row) for row in index_usage]

return metrics

async def get_slow_queries(self) -> List[Dict[str, Any]]:
    """Get slow query statistics"""
    async with self.pool.acquire() as conn:
        # Requires pg_stat_statements extension
        try:

```

```
slow_queries = await conn.fetch("""  
    SELECT query, calls, total_time, mean_time,  
          stddev_time, rows, 100.0 * shared_blk_hit /  
          nullif(shared_blk_hit + shared_blk_read, 0) AS hit_percent  
     FROM pg_stat_statements  
    WHERE mean_time > 100 -- queries taking more than 100ms on average  
    ORDER BY mean_time DESC  
   LIMIT 20  
""")  
    return [dict(row) for row in slow_queries]  
except Exception as e:  
    logger.warning(f"pg_stat_statements not available: {e}")  
    return []  
  
async def check_replication_lag(self) -> Dict[str, Any]:  
    """Check replication lag if applicable"""  
    async with self.pool.acquire() as conn:  
        try:  
            # For primary server  
            lag_info = await conn.fetchrow("""  
                SELECT  
                    client_addr,  
                    state,  
                    pg_wal_lsn_diff(pg_current_wal_lsn(), flush_lsn) as lag_bytes  
               FROM pg_stat_replication  
            """)  
  
            if lag_info:  
                return dict(lag_info)  
            else:  
                return {'status': 'no_replicas'}  
  
        except Exception as e:  
            logger.warning(f"Error checking replication: {e}")  
            return {'error': str(e)}  
  
    async def generate_report(self) -> Dict[str, Any]:  
        """Generate comprehensive monitoring report"""  
        timestamp = datetime.utcnow().isoformat()  
  
        metrics = await self.get_database_metrics()  
        slow_queries = await self.get_slow_queries()  
        replication_info = await self.check_replication_lag()  
  
        report = {  
            'timestamp': timestamp,  
            'database_metrics': metrics,
```

```
'slow_queries': slow_queries,
'replication_info': replication_info,
'health_status': 'healthy' # Default, can be updated based on thresholds
}

# Determine health status
if metrics['connections'] > 80:
    report['health_status'] = 'warning'
if metrics['cache_hit_ratio'] < 95:
    report['health_status'] = 'warning'
if metrics['long_queries'] > 5:
    report['health_status'] = 'critical'

return report

async def close_connection(self):
    """Close database connection pool"""
    if self.pool:
        await self.pool.close()

async def main():
    """Main monitoring loop"""
    postgres_config = {
        'host': 'myapp-postgresql.postgres.database.azure.com',
        'port': 5432,
        'user': 'pgadmin',
        'password': 'SecurePassword123!',
        'database': 'myapp_database'
    }

    monitor = PostgreSQLMonitor(postgres_config)
    await monitor.init_connection()

    try:
        while True:
            report = await monitor.generate_report()

            # Log report
            logger.info(f"Database Health: {report['health_status']}")
            logger.info(f"Active Connections: {report['database_metrics']['connections']}")
            logger.info(f"Cache Hit Ratio: {report['database_metrics']['cache_hit_ratio']}%")
            logger.info(f"Long Running Queries: {report['database_metrics']['long_queries']}")

            # Save report to file or send to monitoring system
            with open(f"/tmp/db_report_{int(time.time())}.json", 'w') as f:
                json.dump(report, f, indent=2, default=str)
```

```
# Send alerts if needed
if report['health_status'] in ['warning', 'critical']:
    logger.warning(f"Database health issue detected: {report['health_status']}")

    # Here you would integrate with your alerting system
    # await send_alert(report)

    # Wait 5 minutes before next check
    await asyncio.sleep(300)

except KeyboardInterrupt:
    logger.info("Monitoring stopped by user")
finally:
    await monitor.close_connection()

if __name__ == "__main__":
    asyncio.run(main())
```

## Cleanup and Optimization

### 9.1 Post-Migration Cleanup

```
bash
```

```
#!/bin/bash
# cleanup_script.sh

echo "Post-Migration Cleanup and Optimization"
echo "====="

# 1. Optimize PostgreSQL configuration
echo "1. Optimizing PostgreSQL configuration..."

# Update PostgreSQL parameters for production workload
az postgres flexible-server parameter set \
--resource-group mysql-to-pg-migration \
--server-name myapp-postgresql \
--name shared_preload_libraries \
--value 'pg_stat_statements,auto_explain'

az postgres flexible-server parameter set \
--resource-group mysql-to-pg-migration \
--server-name myapp-postgresql \
--name log_statement \
--value 'mod'

az postgres flexible-server parameter set \
--resource-group mysql-to-pg-migration \
--server-name myapp-postgresql \
--name log_min_duration_statement \
--value '1000'

# 2. Create indexes for optimal performance
echo "2. Creating optimized indexes..."
PGPASSWORD='SecurePassword123!' psql \
-h myapp-postgresql.postgres.database.azure.com \
-U pgadmin \
-d myapp_database \
-c "
-- Analyze all tables to update statistics
ANALYZE;

-- Create composite indexes based on query patterns
CREATE INDEX CONCURRENTLY IF NOT EXISTS idx_users_email_active
ON users(email) WHERE active = true;

CREATE INDEX CONCURRENTLY IF NOT EXISTS idx_orders_user_date
ON orders(user_id, created_at DESC);

CREATE INDEX CONCURRENTLY IF NOT EXISTS idx_order_items_order_product
```

```

ON order_items(order_id, product_id);

-- Vacuum and analyze after index creation
VACUUM ANALYZE;
"

# 3. Setup automated maintenance
echo "3. Setting up automated maintenance..."
cat > /tmp/pg_maintenance.sql << EOF
-- Create maintenance procedures
CREATE OR REPLACE FUNCTION automated_maintenance()
RETURNS void AS $$
BEGIN
    -- Update table statistics
    ANALYZE;

    -- Vacuum tables with high update/delete activity
    VACUUM (ANALYZE, VERBOSE) users;
    VACUUM (ANALYZE, VERBOSE) orders;
    VACUUM (ANALYZE, VERBOSE) order_items;

    -- Reindex if needed
    REINDEX INDEX CONCURRENTLY idx_users_email_active;

    RAISE NOTICE 'Maintenance completed at %', now();
END;
$$ LANGUAGE plpgsql;

-- Schedule maintenance (requires pg_cron extension)
-- SELECT cron.schedule('pg-maintenance', '0 2 * * 0', 'SELECT automated_maintenance();');

EOF

PGPASSWORD='SecurePassword123!' psql \
-h myapp-postgresql.postgres.database.azure.com \
-U pgadmin \
-d myapp_database \
-f /tmp/pg_maintenance.sql

# 4. Clean up migration resources
echo "4. Cleaning up migration resources..."

# Stop AWS DMS task
aws dms stop-replication-task \
--replication-task-arn arn:aws:dms:us-east-1:account:task:mysql-to-postgresql-continuous

# Delete AWS DMS resources (run after confirming stability)
read -p "Delete AWS DMS resources? (y/N): " confirm

```

```

if [[ $confirm =~ ^[Yy]$ ]]; then
    aws dms delete-replication-task \
        --replication-task-arn arn:aws:dms:us-east-1:account:task:mysql-to-postgresql-continuous

    aws dms delete-endpoint \
        --endpoint-arn arn:aws:dms:us-east-1:account:endpoint:mysql-source-azure

    aws dms delete-endpoint \
        --endpoint-arn arn:aws:dms:us-east-1:account:endpoint:postgresql-target-azure

    aws dms delete-replication-instance \
        --replication-instance-arn arn:aws:dms:us-east-1:account:rep:mysql-to-postgresql-migration
fi

# 5. Update backup strategy
echo "5. Updating backup strategy..."
az postgres flexible-server backup list \
    --resource-group mysql-to-pg-migration \
    --server-name myapp-postgresql

# Configure backup retention
az postgres flexible-server parameter set \
    --resource-group mysql-to-pg-migration \
    --server-name myapp-postgresql \
    --name backup_retention_days \
    --value 30

echo "Cleanup and optimization completed!"

```

## Summary

This comprehensive migration strategy provides:

- 1. Zero-downtime migration** using AWS DMS with continuous replication
- 2. Fallback options** with Debezium/Kafka for real-time CDC
- 3. Application-level dual-write** strategy for seamless cutover
- 4. Comprehensive validation** at every step
- 5. Automated rollback** procedures
- 6. Production-ready monitoring** and alerting
- 7. Performance optimization** post-migration

## Key Success Factors:

- **AWS DMS** handles cross-engine conversion automatically

- **Dual-write pattern** ensures data consistency during transition
- **Comprehensive testing** validates functionality before cutover
- **Monitoring and alerting** catch issues early
- **Rollback procedures** provide safety net

The entire process typically takes 2-4 weeks including planning, testing, and monitoring phases, but the actual downtime is less than 5 minutes during cutover.